

In the name of women, life, freedom

Fatemeh Salehi – 810198423

### Project 3

## Getting Familiar with the game's environment

**Windy Grid world** is a standard grid world with start and goal states. The difference is that a crosswind runs upward through the middle of the grid. Actions are the standard four: up, right, down, and left. In the middle region, the resultant next states are shifted upward by the "wind" whose strength varies from column to column. The reward is -1 until the goal state is reached.




Fig 1: windy grid world

## Q-Learning

Q-learning is a model-free, off-policy reinforcement learning that will find the best course of action, given the agent's current state. Depending on where the agent is in the environment, it will decide the next action to be taken.

**Step 1:** Create an initial Q-Table with all values initialized to 0 and an epsilon initialized to 1, which will degrade at a constant rate.

**Step 2:** If epsilon is less than 0.05 (or some other threshold) choose a random action and perform it unless choosing the action with the highest value. Update values in the table.

**Step 3:** Get the value of the reward and calculate the value Q-Value using Bellman Equation

**Step 4:** Continue the same until an episode ends

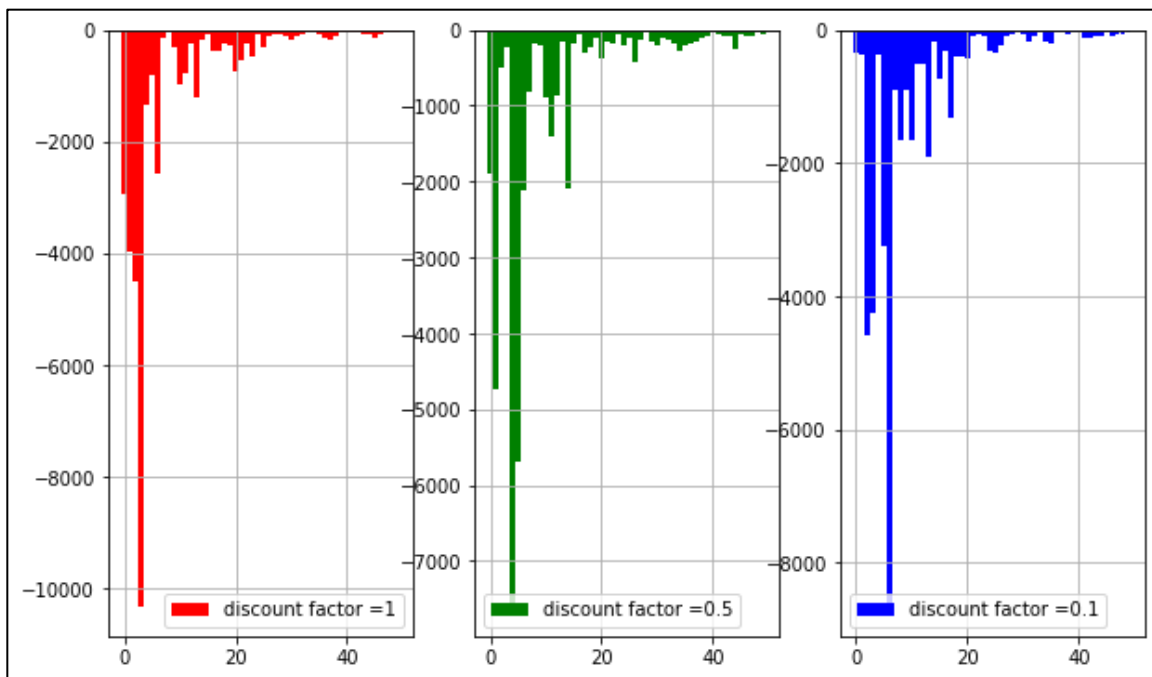
$$\text{New } Q(S, A) = \text{Current } Q \text{ Value} + \text{Learning Rate} \left[ \text{Reward} + \text{Discount Rate} \times \text{Maximum Expected Future Reward} - \text{Current } Q \text{ Value} \right]$$

$\text{New } Q(S, A) = Q(S, A) + \alpha [R(S, A) + \gamma \text{Max } Q'(S', A') - Q(S, A)]$

Current Q Value    Learning Rate    Reward    Discount Rate    Maximum Expected Future Reward

**Fig 2: Bellman Equation**

We ran the code for 3 different values of the discount factor and the results are as follows:



**Fig 3: the rewards in each episode of Q-learning for different discount values**

According to Fig 3, the greater the discount factor, the faster the convergence rate; it is worth mentioning that for smaller discount factors, the Q-values are stuck in a local optimal solution, which makes the overall reward not change over a period.

## Actor-Critic

In this algorithm, the temporal difference is defined by value function

$$\delta_t = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

$V(s_{t+1}) \rightarrow$  The value of the next state may give more insight about how the last action was good rather than blending this value to the action taken in the next step.

$$\text{softmax policy: } \pi_t(s, a) = \Pr\{a_t = a | s_t = s\} = \frac{e^{P(s,a)}}{\sum_b e^{P(s,b)}}$$

$$P(s_t, a_t) \leftarrow P(s_t, a_t) + \beta \delta_t$$

$$V_{t+1}(s_t) = V_t(s_t) + \alpha \delta_t = V_t(s_t) + \alpha [r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

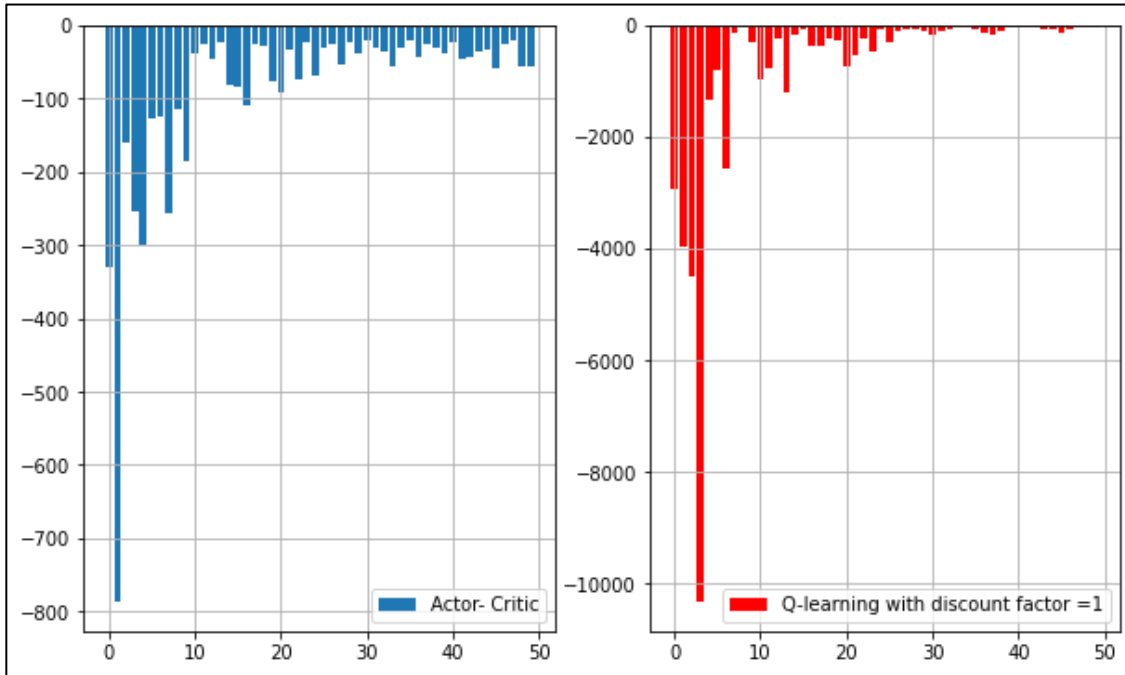



Fig 4: The rewards in each episode of Q-learning with a discount factor of 1 and Actor-Critic

According to Fig 4, we can conclude that Q-learning works better than Actor-Critic, which means that the action that the agent takes in the next state is critical, and only state value is not enough to evaluate a state.

## DDQN (Double Deep Q-Network)

This method involves using two separate Q-value estimators, each of which is used to update the other. Using these independent estimators, we can unbiased Q-value estimates of the actions selected using the opposite estimator [1]. We can thus avoid maximization bias by disentangling our updates from biased estimates.

### Algorithm 1 : Double Q-learning (Hasselt et al., 2015)

```
Initialize primary network  $Q_\theta$ , target network  $Q_{\theta'}$ , replay buffer  $\mathcal{D}$ ,  $\tau < 1$ 
for each iteration do
  for each environment step do
    Observe state  $s_t$  and select  $a_t \sim \pi(a_t, s_t)$ 
    Execute  $a_t$  and observe next state  $s_{t+1}$  and reward  $r_t = R(s_t, a_t)$ 
    Store  $(s_t, a_t, r_t, s_{t+1})$  in replay buffer  $\mathcal{D}$ 
  for each update step do
    sample  $e_t = (s_t, a_t, r_t, s_{t+1}) \sim \mathcal{D}$ 
    Compute target Q value:
       $Q^*(s_t, a_t) \approx r_t + \gamma Q_{\theta'}(s_{t+1}, \operatorname{argmax}_{a'} Q_{\theta'}(s_{t+1}, a'))$ 
    Perform gradient descent step on  $(Q^*(s_t, a_t) - Q_\theta(s_t, a_t))^2$ 
    Update target network parameters:
       $\theta' \leftarrow \tau * \theta + (1 - \tau) * \theta'$ 
```

The output of this algorithm for batch-size of 64 is as following:

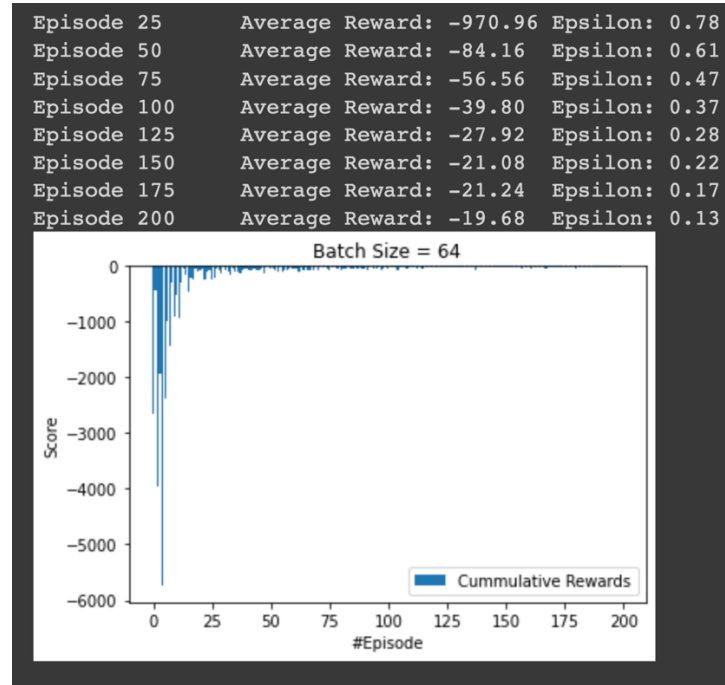


Fig 5: The rewards in each episode of DDQN with batch size = 64

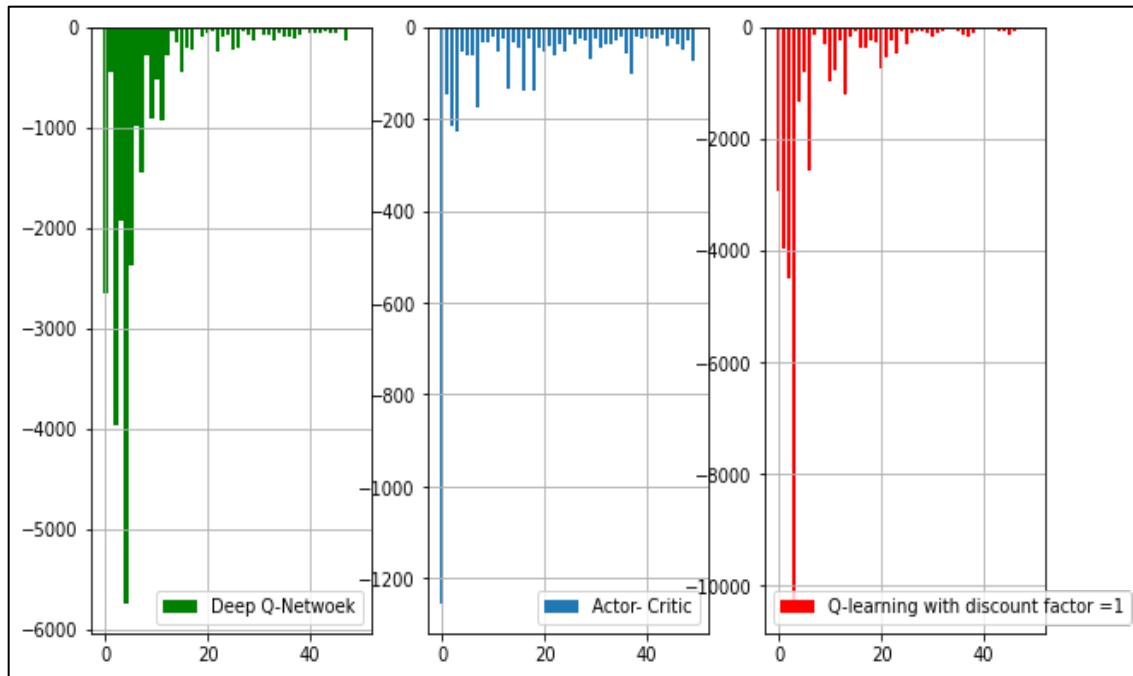


Fig 6: The rewards in each episode of Q-learning with a discount factor of 1, Actor-Critic, and DDQN

According to Fig 6, DDQN converges sooner than both Q-learning and Actor-Critic even though it is more complex than the other two.

## Compare the taken paths

s = (3, 0)	s = (3, 0)	s = (3, 0)
Right	Right	Right
s = (3, 1)	s = (3, 1)	s = (3, 1)
Right	Right	Right
s = (3, 2)	s = (3, 2)	s = (3, 2)
Right	Right	Right
s = (3, 3)	s = (3, 3)	s = (3, 3)
Right	Right	Right
s = (2, 4)	s = (2, 4)	s = (2, 4)
Right	Right	Right
s = (1, 5)	s = (1, 5)	s = (1, 5)
Right	Right	Right
s = (0, 6)	s = (0, 6)	s = (0, 6)
Right	Right	Right
s = (0, 7)	s = (0, 7)	s = (0, 7)
Right	Right	Right
s = (0, 8)	s = (0, 8)	s = (0, 8)
Right	Right	Right
s = (0, 9)	s = (0, 9)	s = (0, 9)
Up	Up	Up
s = (1, 9)	s = (1, 9)	s = (1, 9)
Up	Up	Up
s = (2, 9)	s = (2, 9)	s = (2, 9)
Up	Up	Up
s = (3, 9)	s = (3, 9)	s = (3, 9)
Up	Up	Up
s = (4, 9)	s = (4, 9)	s = (4, 9)
Left	Left	Left
s = (4, 8)	s = (4, 8)	s = (4, 8)
Left	Left	Left
s = (3, 7)	s = (3, 7)	s = (3, 7)

Fig 6: The paths taken by Q-learning, Actor-Critic, and DDQN respectively

## Multi-agent RL for two agents in the windy grid world

In the environment, two agents will search for a way to the goal position without bumping into each other. The reward for each action is -1; if the agents hit each other the reward would be -10, and if agents are both at the green position, they get the reward of 10.

### The Friend-or-Feo Q-Learning


In windy grid world, the agents are friends meaning that they want to help each other out in order to get to the goal.

$$V_i(s) \leftarrow \max_{\substack{a_1 \in A_1 \\ a_2 \in A_2}} Q_i(s, a_1, a_2) \quad i = 1, 2$$

$$Q_i(s, a_1, a_2) \leftarrow Q_i(s, a_1, a_2) + \alpha[r + \gamma V_i(s') - Q_i(s, a_1, a_2)] \quad i = 1, 2$$

#### Algorithm 4 Nash-Q for player $k$

**Input:** learning rate  $\alpha$ , discount factor  $\gamma$   
**for all**  $s \in S, i \in K, a^1 \in A^1, \dots, a^n \in A^n$  **do**  
     $Q^i(s, a^1, \dots, a^n) \leftarrow 0$   
     $(\pi^1(s, \cdot), \dots, \pi^n(s, \cdot)) \leftarrow \text{Nash}(Q^1(s), \dots, Q^n(s))$   
     $V^i(s) \leftarrow \sum_{a^1, \dots, a^n} Q^i(s, a^1, \dots, a^n) \prod_{j=1}^n \pi^j(s, a^j)$   
**end for**  
Observe the current state  $s$   
**loop**  
    Choose action  $a^k$  for state  $s$  using policy  $\pi^k$  (with proper exploration)  
    Take action  $a^k$ , observe other agents' actions  $a^1, \dots, a^{k-1}, a^{k+1}, \dots, a^n$ , reward  $r^k$ , other agents' rewards  $r^1, \dots, r^{k-1}, r^{k+1}, \dots, r^n$  and succeeding state  $s'$  provided by the environment  
    **for all**  $i \in K$  **do**  
         $Q^i(s, a^1, \dots, a^n) \leftarrow Q^i(s, a^1, \dots, a^n) + \alpha [r^i + \gamma V^i(s') - Q^i(s, a^1, \dots, a^n)]$   
    **end for**  
     $(\pi^1(s, \cdot), \dots, \pi^n(s, \cdot)) \leftarrow \text{Nash}(Q^1(s), \dots, Q^n(s))$   
    **for all**  $i \in K$  **do**  
         $V^i(s) \leftarrow \sum_{a^1, \dots, a^n} Q^i(s, a^1, \dots, a^n) \prod_{j=1}^n \pi^j(s, a^j)$   
    **end for**  
    decay  $\alpha$   
     $s \leftarrow s'$   
**end loop**



**Fig 7: The rewards in each episode of Friend-or-Feo method**


## The Distributed Q-learning

The idea is that each agent ignores the fact that other agents exist and only consider its local Q-table. The local Q-values are updated only when the update leads to an increase in the Q-value.

$$Q_{i,k+1}(x_k, u_{i,k}) = \max \{Q_{i,k}(x_k, u_{i,k}), r_{k+1} + \gamma \max_{u_i} Q_{i,k}(x_{k+1}, u_{i,k})\}$$

The local policy is updated only if the update leads to an improvement in the Q-value.

$$\bar{h}_{i,k+1}(x_k) = \begin{cases} u_{i,k} & \text{if } \max_{u_i} Q_{i,k+1}(x_k, u_i) > \max_{u_i} Q_{i,k}(x_k, u_i) \\ \bar{h}_{i,k}(x_k) & \text{otherwise} \end{cases}$$



**Fig 8: The rewards in each episode of Distributed Q-learning method**

## Fully Cooperative Method

The agents have the same reward function and the same Q-table; the learning goal is to maximize the common discount return.

$$Q_{k+1}(x_k, u_k) = Q_k(x_k, u_k) + \alpha \left[ r_{k+1} + \gamma \max_{u'} Q_k(x_{k+1}, u') - Q_k(x_k, u_k) \right]$$




Fig 9: The rewards in each episode of Fully Cooperative method

### Conclusion:

According to the Fig7, 8, and 9 the best method for this environment is Distributed Q-learning because it converges way sooner than the other methods.

References:

- [1] “Deep Reinforcement Learning with Double Q-learning” (Hasselt et al., 2015).